

Introduction to Bioinformatics

Short Read Alignment

Ali Sharifi-Zarchi Somayyeh Koohi

CE Department
Sharif University of Technology

Fall 2025-26

Presentation Overview

- 1 Problem & Challenges
- 2 Brute-Force: Pairwise Alignment with DP
 - Dynamic Programming Foundations
 - Types of Sequence Alignment
 - Global Alignment - Needleman-Wunsch
 - Local Alignment - Smith-Waterman
 - Affine Gaps (Practical Improvement)
- 3 Why Pure DP is Too Slow for Read Alignment
- 4 Burrows-Wheeler Transform & FM-Index
 - BWT Construction Example
 - Summary
- 5 References

The Read Alignment Problem

- DNA sequencing → millions of short reads (50–150 bp)
- We have a long reference genome (e.g., human: 3 billion bp)
- **Goal:** For each read, find where (if anywhere) it came from in the reference

⇒ Massive approximate string matching

Real-World Challenges

- Reference size: bacteria $\sim 5\text{Mbp}$, human $\sim 3\text{Gbp}$
- Number of reads: $10^6 - 10^9$
- Reads are short and noisy (sequencing errors, SNPs, indels)
- Genomes are highly repetitive
- Must finish in hours on limited RAM/CPU

What Do Modern Aligners Need?

Key Requirements

- High sensitivity (find true origin even with a few errors)
- Very high speed & scalability
- Low memory usage (index must fit in RAM)
- Reference indexed once, queried millions of times

Pairwise Sequence Alignment


Core idea

Compare a short read (pattern) against every possible position in the reference, allowing mismatches and gaps.

Two classic formulations:

- **Global alignment** (Needleman-Wunsch) - entire read
- **Local alignment** (Smith-Waterman) - best substring

The Manhattan Tourist Problem (warm-up)



Find the maximum-weight path from
top-left to bottom-right



Recurrence:

$$s_{i,j} = \max \begin{cases} s_{i-1,j} + w(\text{vertical}) \\ s_{i,j-1} + w(\text{horizontal}) \end{cases}$$

General DP Strategy

- 1 Turn the problem into a path-finding problem in a grid
- 2 Define a recurrence: value of cell (i, j) depends only on a few previous cells
- 3 Fill table row-by-row (or column-by-column)
- 4 Trace back to recover the actual alignment

Global, Local, Multiple

Target Sequence

5' ACTACTAGATTACTTACGGATCAGGTACTTTAGAGGCTTGCAACCA 3'

Query Sequence

Query Sequence 5' TACTCACGGATGAGGTACTTTAGAGGC 3'

Global Alignment

Target Sequence

5' ACTACTAGATTACTTACGGATCAGGTACTTTAGAGGCTTGCAACCA 3'

5' ACTACTAGATT---ACGGATC--GTACTTTAGAGGCTAGCAACCA 3'

Query Sequence

Multiple Sequence Alignment (MSA)



Figure: Comparison between the kinds of alignment

Needleman-Wunsch (1970) - Global Alignment

Recurrence (simplified scoring: match +1, mismatch/gap -1):

$$F(i, j) = \max \begin{cases} F(i-1, j-1) + s(A_i, B_j) & \text{diagonal (match/mismatch)} \\ F(i-1, j) - d & \text{vertical (gap in B)} \\ F(i, j-1) - d & \text{horizontal (gap in A)} \end{cases}$$

Needleman-Wunsch Example

Scoring: match = +1, mismatch = -1, gap = -1

	-	A	C	T	G	A	T	C	A
-	0	-1	-2	-3	-4	-5	-6	-7	-8
A	-1	1	0	-1	-2	-3	-4	-5	-6
T	-2	0	0	1	0	-1	-2	-3	-4
G	-3	-1	-1	0	2	1	0	-1	-2
A	-4	-2	-2	-1	1	3	2	1	0
T	-5	-3	-3	-1	0	2	4	3	2
C	-6	-4	-3	-2	-1	1	3	5	4
A	-7	-5	-4	-3	-2	0	2	4	6

Final score = 6 $\Rightarrow$ very good global alignment

- **Red 0** = starting point (empty alignment)

Smith-Waterman (1981) - Local Alignment

Same three choices, but we are allowed to start fresh:

$$H(i, j) = \max \begin{cases} 0 \\ H(i-1, j-1) + s(A_i, B_j) \\ H(i-1, j) - d \\ H(i, j-1) - d \end{cases}$$

- 0 means "start a new alignment here"
- Optimal local alignment = maximum value in the entire matrix
- Trace back from that maximum until reaching 0

Global vs. Local


```
--T--CC-C-AGT--TATGT-CAGGGGACACG--A-GCATGCAGA-GAC
  |  || |  ||  | | |  || |  | |  ||||  |
AATTGCCGCC-GTCGT-T-TTCAG----CA-GTTATG--T-CAGAT--C
```

Figure: Global alignment example

```
          tccCAGTTATGTCAGgggacacgagcatgcagagac
          |||||
aattgccgccgtcgttttcagCAGTTATGTCAGatc
```

Figure: Local alignment example

Global vs. Local



Affine Gap Penalties

Real gaps are usually long $\rightarrow$ opening a gap is more expensive than extending it.

$\Rightarrow$ Use three matrices (Gotoh 1982):

- $M(i, j)$ - best alignment ending with match/mismatch
- $I_x(i, j)$ - best ending with gap in sequence x
- $I_y(i, j)$ - best ending with gap in sequence y

Widely used in BLAST, BWA, Bowtie2, etc.

Complexity Problem

- Needleman-Wunsch / Smith-Waterman: $O(mn)$ time and space
- m = read length (100), n = reference length (3×10^9)
- Trying every position $\rightarrow 3 \times 10^{11}$ cell computations per read
- Even at 10^9 cells/sec $\rightarrow$ hours per read $\rightarrow$ impossible

Conclusion: We cannot afford full DP for every read and every position.

Triple Sequence Alignment

Example

```
--T--CC-C-AGT--TATGT-CAGGGGACACG--A-GCATGCAGA-GAC
      |  || |  ||  |  |  |  ||  ||  |  |  |  ||||  |
AATTGCCGCC-GTCGT-T-TTCAG-----CA-GTTATG--T-CAGAT--C
      |||||  |  X|||  |  ||  XXX|||  |  |||  |  |
-ATTGC-G--ATTCGTAT-----GGGACA-TGGATGCATGCAG-TGAC
```

Figure: Triple alignment example



Triple Sequence Alignment

Policy

$$s_{i,j,k} = \max \left\{ \begin{array}{ll} s_{i-1,j,k} & +\delta(v_i, -, -) \\ s_{i,j-1,k} & +\delta(-, w_j, -) \\ s_{i,j,k-1} & +\delta(-, -, u_k) \\ s_{i-1,j-1,k} & +\delta(v_i, w_j, -) \\ s_{i-1,j,k-1} & +\delta(v_i, -, u_k) \\ s_{i,j-1,k-1} & +\delta(-, w_j, u_k) \\ s_{i-1,j-1,k-1} & +\delta(v_i, w_j, u_k) \end{array} \right.$$

Triple Sequence Alignment

3-dimensional DP table



Summary & Transition

What we learned

- Pairwise alignment → solved elegantly with DP
- Global (Needleman-Wunsch) and local (Smith-Waterman)
- Affine gaps make scoring biologically realistic

But

Pure DP is far **too slow** for millions of reads against huge genomes.

Why the Burrows-Wheeler Transform (BWT)?

Real-world impact

- Illumina NovaSeq: > 3 terabases of DNA per day
- Human genome index must fit in RAM and answer queries in microseconds

BWT powers nearly every modern short-read aligner:

- Bowtie / Bowtie2 BWA (MEM & ALN) minimap2 SNAP GEM ...

⇒ The single most important data structure in bioinformatics today

BWT Construction – “panama\$” Example

$S = \text{panama\$}$ (\$ is smallest sentinel)

#	Rotation	Sorted rotation	BWT
0	panama\$	\$panama	a
1	anama\$p	a\$panam	m
2	nama\$pa	ama\$pan	n
3	ama\$pan	anama\$p	p
4	ma\$pana	ma\$pana	a
5	a\$panam	nama\$pa	a
6	\$panama	panama\$	\$

$\Rightarrow$ **BWT = amnpaa\$**

The Full BWT Matrix with Ranks (LF Mapping)

BWT = amnpaa\$

i	F (sorted)	L[i]	Rank in L	Rank in F $\rightarrow$ LF(i)
0	\$panama	a	a^1	$a^1 \rightarrow 1$
1	a\$panam	m	m^1	$m^1 \rightarrow 4$
2	ama\$pan	n	n^1	$n^1 \rightarrow 5$
3	anama\$p	p	p^1	$p^1 \rightarrow 6$
4	ma\$pana	a	a^2	$a^2 \rightarrow 2$
5	nama\$pa	a	a^3	$a^3 \rightarrow 3$
6	panama\$	\$	$\1	$\$^1 \rightarrow 0$

LF rule: k -th c in $L \rightarrow k$ -th c in F

FM-Index: What We Store in Practice

C[c] = count of chars $< c$

- $C[\$]=0, C[a]=1, C[m]=4, C[n]=5, C[p]=6$

Occ(c,i) = occurrences of c in $BWT[0..i]$

LF(i) = C[BWT[i]] + Occ(BWT[i], i-1)

$\Rightarrow$ Human genome index $\approx 3-5$ GB

Constructing BWT from Suffix Array

Python Code

```
def compute_bwt(s, sa):  
    bwt = []  
    for p in sa:  
        if p == 0:  
            bwt.append('$')  
        else:  
            bwt.append(s[p-1])  
    return ''.join(bwt)
```

Caveat

Constructing the full suffix array still costs $O(n \log n)$ time and $O(n)$ extra space $\Rightarrow$ this removes the memory benefits of the BWT!

The Crucial Property: Last-to-First (LF) Mapping

Index i	$L[i]$ (BWT)	$F[i]$ (first column)
0	t	\$
1	t	a
2	t	a
3	t	a
4	\$	c
5	a	t
6	a	t
7	a	t
8	c	t

LF mapping rule: If $L[i] = c$ is the k -th occurrence of c in the BWT (column L), then $LF(i)$ is the row where the k -th occurrence of c appears in column F .

⇒ Following LF repeatedly walks **backwards** through the original string!

Decoding the Original String using LF Mapping

Python code

```
def decode_bwt(bwt, LF):
    n = len(bwt)
    r = 0                # start at row containing $
    s = ['$']
    while len(s) < n:
        s.append(bwt[r])  # prepend the previous character
        r = LF[r]         # go to previous position
    return ''.join(reversed(s))
```

Result: Starting from the \$ row and repeatedly applying LF reconstructs the original text in reverse.

LF Mapping in Practice – C and Occ arrays

We don't store the full LF array (would cost $O(n)$ space). Instead we use two small tables:

- $C[c]$ = number of characters **strictly smaller** than c in the text
- $\text{Occ}(c, i)$ = number of occurrences of c in $\text{BWT}[0..i]$

Then:

$$\text{LF}(i) = C[\text{BWT}[i]] + \text{Occ}(\text{BWT}[i], i - 1)$$

⇒ This is the core of the **FM-index** (Ferragina–Manzini 2000).

FM-Index = Compressed Full-Text Index

The FM-index supports **backward searching** of patterns in $O(m)$ time (often sub-linear with sampling):

- 1 Start with interval $[0, n)$
- 2 For each character c of the pattern **from right to left**:

$$i' = C[c] + \text{Occ}(c, i - 1)$$

$$j' = C[c] + \text{Occ}(c, j) - 1$$

- 3 If $i' \leq j' \Rightarrow$ continue, else pattern not present

This is exactly how Bowtie, BWA-MEM, etc. find seed locations fast.

Backward Search – Find “ana” in “panama\$”

Start: [0,7) Pattern: $a \rightarrow n \rightarrow a$ (right to left)

Step	Char	Calculation	Interval
0	-	-	[0,7)
1	a	$1+0=1 \dots 1+3=4$	alert[1,4)
2	n	$5+0=5 \dots 5+1=6$	alert[5,6)
3	a	$1+2=3 \dots 1+3=4$	alert[3,4) Success

“ana” occurs once, starting at position 2

Backward Search – Visualizing Step 1 (final “a”)

Current interval $[0,7)$ $\rightarrow$ looking for preceding **a**

i	BWT[i]	F row
0	a	\$panama
1	m	a\$panam $\rightarrow a^1$
2	n	ama\$pan $\rightarrow a^2$
3	p	anama\$p $\rightarrow a^3$
4	a	ma\$pana
5	a	nama\$pa
6	\$	panama\$

$3 \times a \rightarrow$ new interval $[1,4)$

Backward Search – Step 2: Process “n”

Current interval: $[1,4)$ → rows 1–3 only

We are now looking for preceding character **n**

Only row 2 has $\text{BWT}[2] = \text{n}$ in current range → its preceding char (via LF) is n

The 1st ‘n’ in F is at row 5 → new interval becomes $[5,6)$

Backward Search – Step 3: Process first “a” again

Current interval: $[5,6)$ → only row 5
BWT[5] = a → matches our pattern ‘a’
→ maps to row 3 (3rd ‘a’ in F)

Final interval $[3,4)$ → exact match found!

Conclusion: “ana” occurs once, starting at position 2 in “panama\$”

Summary – Why BWT is Perfect for Short-Read Alignment

- Reversible transformation
- Highly compressible (runs of identical characters)
- Enables **exact** backward search in $O(m)$ time using only $O(n)$ bits
- With sampling $\Rightarrow$ locate occurrences in constant time
- Whole human genome index fits in ≤ 4 **GB** RAM
- Used by virtually every modern short-read aligner (Bowtie, BWA, etc.)

Next step: Seed-and-extend using these exact matches $\rightarrow$ fast, sensitive alignment.

Why is Seed & Extend Necessary?

- A full exact BWT search works well only when the read has **no errors**.
- Real sequencing data contains:
 - Substitutions
 - Insertions and deletions
 - PCR/sequencing errors
- Supporting errors directly inside the BWT search is expensive.
- **Solution:** Modern aligners index the genome but search only short exact (or near-exact) **seeds**.

Key Idea

Instead of aligning the whole read, find a reliable seed quickly, and then extend it locally to reconstruct the full alignment.

- This reduces the search space by orders of magnitude.
- Makes BWT-based alignment practical, fast, and scalable.

Full Alignment vs. Seed & Extend

Full Alignment (DP or full BWT search)

- Attempts to align the entire read at once.
- Expensive in both time and memory.
- Error tolerance requires exploring many paths.
- Not scalable for billions of bases.

Seed & Extend

- Search only short seeds in the index.
- Very fast FM-index exact matching on seeds.
- Extension uses a small DP or local BWT search.
- Scales to whole human genome efficiently.

Conclusion

Seed & Extend combines:

- **Speed** from FM-index seed lookup, and
- **Accuracy** from local DP extension.

This is why all modern short-read aligners use it.

Sina Beyrami
Sina Daneshgar

References

- **Istrail, S., Pevzner, P.** *An Introduction to Bioinformatics Algorithms.*
- **Needleman and Wunsch** (1970) *A general method applicable to the search for similarities in the amino acid sequence of two proteins*
- **Smith and Waterman** (1981) *Identification of common molecular subsequences*
- **Ugene** Alignment of Short Reads
- **Wikipedia**
- **Sven Rahmann** (2021). *The Burrows-Wheeler Transform (BWT) – Algorithms for Sequence Analysis.*
- **Burrows M., Wheeler D. J.** (1994). *A block-sorting lossless data compression algorithm.*
- **Ferragina P., Manzini G.** (2000). *Opportunistic data structures with applications.*
- **Simpson J. T., Durbin R.** (2010). *Efficient construction of an assembly string graph using the FM-index.*
- **Langmead B., Trapnell C., Pop M., Salzberg S. L.** (2009). *Ultrafast and memory-efficient alignment of short DNA sequences to the human genome.*
- **Li H., Durbin R.** (2009). *Fast and accurate short read alignment with Burrows-Wheeler transform.*